

# Time Series

how you mark and refer to time series data depends on the application, and you may have one of the following

- Timestamps, specific instants in time
- fixed periods, such as the month jan 2007 or the full year 2010
- intervals of time, indicated by a start and end timestamp. Periods can be thought of as special cases of intervals
- Experiment or elapsed time; each timestamp is a measure of time relative to a particular start time(e.g. the diameter of a cookie baking each second since being placed in the oven)

```
In [44]: import pandas as pd
```

```
In [45]: from datetime import datetime
now = datetime.now()
now
```

```
Out[45]: datetime.datetime(2020, 8, 4, 15, 45, 15, 197133)
```

```
In [46]: now.year, now.month, now.day
```

```
Out[46]: (2020, 8, 4)
```

```
In [47]: delta = datetime(2011, 1, 7) - datetime(2008, 6, 24, 8, 15)
delta
```

```
Out[47]: datetime.timedelta(926, 56700)
```

```
In [48]: from datetime import timedelta
start = datetime(2011, 1, 7)
start + 2 * timedelta(12)
```

```
Out[48]: datetime.datetime(2011, 1, 31, 0, 0)
```

## Converting Between String and Datetime

```
In [49]: stamp = datetime(1998, 12, 28)
str(stamp)
```

```
Out[49]: '1998-12-28 00:00:00'
```

```
In [50]: stamp.strftime('%Y-%m-%d')
```

```
Out[50]: '1998-12-28'
```

## Datetime format specification

%Y -- four

%y -- Two-digit year

%m -- Two-digit month

%d -- Two-digit day

%H -- Hour (24 hour clock)

%I -- House (12 hour clock)

%M -- Two-digit minute

%S -- Second [00,61] 61 account for leap second

%w -- Weekday as integer [0(sunday), 6]

%U -- Week number of the year [0-53]; sunday is considered the first day of the week, and days before the first sunday of the year are "week 0"

%W -- Week number of the year [00, 53] Monday is considered the first day of the week, and days before the first monday of the year are week "0"

%z -- UTC time zone offset as +HHMM or -HHMM; empty if time zone naive

%F -- Shortcut for %Y-%m-%d

%D -- Shortcut for %m/%d/%y

you can use the same methods to convert strings to dates using `datetime.strptime`

```
In [51]: value = '2011-01-03'
datetime.strptime(value, '%Y-%m-%d')
```

```
Out[51]: datetime.datetime(2011, 1, 3, 0, 0)
```

```
In [52]: import numpy as np
dates = [datetime(2011, 1, 2), datetime(2011, 1, 5), datetime(2011, 1, 8),
ts = pd.Series(np.random.randn(6), index=dates)
ts
```

```
Out[52]: 2011-01-02    -0.596279
2011-01-05    -0.815065
2011-01-07     1.743168
2011-01-08    -0.773380
2011-01-10    -0.306357
2011-01-12     1.100062
dtype: float64
```

```
In [53]: ts.index
```

```
Out[53]: DatetimeIndex(['2011-01-02', '2011-01-05', '2011-01-07', '2011-01-08',  
                        '2011-01-10', '2011-01-12'],  
                      dtype='datetime64[ns]', freq=None)
```

```
In [54]: ts + ts[::2] # selects every other point and doubles them.
```

```
Out[54]: 2011-01-02    -1.192558  
         2011-01-05         NaN  
         2011-01-07     3.486336  
         2011-01-08         NaN  
         2011-01-10    -0.612714  
         2011-01-12         NaN  
         dtype: float64
```

```
In [55]: stamp = ts.index[0]  
         stamp
```

```
Out[55]: Timestamp('2011-01-02 00:00:00')
```

```
In [56]: ts[stamp]
```

```
Out[56]: -0.5962792122800337
```

```
In [57]: ts['2011/01/02'] # can pass a date as a string
```

```
Out[57]: -0.5962792122800337
```

```
In [58]: #for larger series sets  
         longer_ts = pd.Series(np.random.randn(1000), index=pd.date_range('1/1/  
         longer_ts.tail())
```

```
Out[58]: 2002-09-22    -0.493384  
         2002-09-23    -0.207827  
         2002-09-24    -0.886046  
         2002-09-25     0.237732  
         2002-09-26     1.460096  
         Freq: D, dtype: float64
```

```
In [59]: longer_ts['2001'].tail()
```

```
Out[59]: 2001-12-27    -0.744843  
         2001-12-28    -0.074248  
         2001-12-29    -0.944912  
         2001-12-30     1.131164  
         2001-12-31    -0.408839  
         Freq: D, dtype: float64
```

```
In [60]: longer_ts['2002-05'].tail()
```

```
Out[60]: 2002-05-27    0.718844
          2002-05-28    1.360635
          2002-05-29   -2.496779
          2002-05-30    0.412370
          2002-05-31    1.176583
          Freq: D, dtype: float64
```

```
In [61]: ts[datetime(2009, 5, 5):]
```

```
Out[61]: 2011-01-02   -0.596279
          2011-01-05   -0.815065
          2011-01-07    1.743168
          2011-01-08   -0.773380
          2011-01-10   -0.306357
          2011-01-12    1.100062
          dtype: float64
```

```
In [62]: ts['1/6/2011':'1/11/2011']
```

```
Out[62]: 2011-01-07    1.743168
          2011-01-08   -0.773380
          2011-01-10   -0.306357
          dtype: float64
```

```
In [63]: ts.truncate(after='1/9/2011')
```

```
Out[63]: 2011-01-02   -0.596279
          2011-01-05   -0.815065
          2011-01-07    1.743168
          2011-01-08   -0.773380
          dtype: float64
```

## Time Series with Duplicate Indices

```
In [64]: dates = pd.DatetimeIndex(['1/1/2002', '1/2/2002', '1/3/2002', '1/2/2002'])
          dup_ts = pd.Series(np.random.randn(5), index=dates)
          dup_ts
```

```
Out[64]: 2002-01-01   -1.410824
          2002-01-02    0.881724
          2002-01-03    1.002362
          2002-01-02    0.173960
          2002-01-02    0.180727
          dtype: float64
```

```
In [65]: dup_ts.index.unique()
```

```
Out[65]: DatetimeIndex(['2002-01-01', '2002-01-02', '2002-01-03'], dtype='datetime64[ns]', freq=None)
```

```
In [66]: # suppose you wanted to group non-unique dates  
non_unique = dup_ts.groupby(level=0)  
non_unique.mean()
```

```
Out[66]: 2002-01-01    -1.410824  
2002-01-02     0.412137  
2002-01-03     1.002362  
dtype: float64
```

```
In [67]: non_unique.count()
```

```
Out[67]: 2002-01-01     1  
2002-01-02     3  
2002-01-03     1  
dtype: int64
```

## Date Ranges Frequencies and shifting

often you time series may have non-uniform date ranges. For most purposes this is fine but should you wish to have equidistant dates you may wish to use the following

```
In [68]: ts
```

```
Out[68]: 2011-01-02    -0.596279  
2011-01-05    -0.815065  
2011-01-07     1.743168  
2011-01-08    -0.773380  
2011-01-10    -0.306357  
2011-01-12     1.100062  
dtype: float64
```

```
In [69]: resampler = ts.resample('D') #D for daily
resampler.mean()
```

```
Out[69]: 2011-01-02    -0.596279
2011-01-03         NaN
2011-01-04         NaN
2011-01-05    -0.815065
2011-01-06         NaN
2011-01-07     1.743168
2011-01-08    -0.773380
2011-01-09         NaN
2011-01-10    -0.306357
2011-01-11         NaN
2011-01-12     1.100062
Freq: D, dtype: float64
```

```
In [70]: index = pd.date_range('2009-04-01', '2009-05-01')
index
```

```
Out[70]: DatetimeIndex(['2009-04-01', '2009-04-02', '2009-04-03', '2009-04-04',
                        ,
                        '2009-04-05', '2009-04-06', '2009-04-07', '2009-04-08',
                        ,
                        '2009-04-09', '2009-04-10', '2009-04-11', '2009-04-12',
                        ,
                        '2009-04-13', '2009-04-14', '2009-04-15', '2009-04-16',
                        ,
                        '2009-04-17', '2009-04-18', '2009-04-19', '2009-04-20',
                        ,
                        '2009-04-21', '2009-04-22', '2009-04-23', '2009-04-24',
                        ,
                        '2009-04-25', '2009-04-26', '2009-04-27', '2009-04-28',
                        ,
                        '2009-04-29', '2009-04-30', '2009-05-01'],
                        dtype='datetime64[ns]', freq='D')
```

```
In [71]: index2 = pd.date_range(start='1998-12-28', periods=365) #can also usec
# should you have wanted a index that just referenced the last busines
index2 = pd.date_range(start='1998-12-28', end='2020-12-28', freq='BM')
```

## Base time series frequencies

D -- Day

B -- BusinessDay

H -- Hour

T/min -- Minute

S -- Second

L/ms -- milli

U -- Micro

M -- MonthEnd

BM -- BusinessMonthEnd

MS -- MonthBegin

BMS -- BusinessMonthBegin

W-MON, W-TUE -- Week

WOM-1MON/WOM-2MON -- Generates dates in first/second/third/fourth week of the month.

Q-JAN/Q-FEB -- Quarter end anchored on last calendar name of each month

BQ-JAN/BQ-FEB -- Business Quarter End. Quarterly dates anchored on last weekday of each month

QS-JAN/QS-FEB -- Quarterly dates anchored on first calendar day of each month.

BQS-JAN/BQS-FEB -- Quarterly dates anchored on first weekday of each month, for year ending in indicated month.

A-JAN/A-FEB -- Annual dates anchored on last calendar day of given month

BA-JAN/BA-FEB -- Annual dates anchored on last weekday of given month

AS-JAN/AS-FEB -- Annual dates anchored on first day of given month

BAS-JAN/BAS-FEB -- Annual dates anchored on first weekday of given month.

```
In [72]: pd.date_range('2001-11-23 12:56:31', periods=5)
```

```
Out[72]: DatetimeIndex(['2001-11-23 12:56:31', '2001-11-24 12:56:31',
                        '2001-11-25 12:56:31', '2001-11-26 12:56:31',
                        '2001-11-27 12:56:31'],
                        dtype='datetime64[ns]', freq='D')
```

```
In [73]: pd.date_range('2001-11-23 12:56:31', periods=5, normalize=True) #should
```

```
Out[73]: DatetimeIndex(['2001-11-23', '2001-11-24', '2001-11-25', '2001-11-26',
                        '2001-11-27'],
                        dtype='datetime64[ns]', freq='D')
```

## Shifting Data

moving information around by shifting date

```
In [74]: dates = pd.date_range('1998-12-05', periods = 6)
         ts = pd.Series(np.random.randn(6), index=dates)
         ts
```

```
Out[74]: 1998-12-05    0.484211
         1998-12-06   -1.271930
         1998-12-07    0.467360
         1998-12-08    0.952550
         1998-12-09    0.371530
         1998-12-10   -1.454242
         Freq: D, dtype: float64
```

```
In [75]: ts.shift(3)
```

```
Out[75]: 1998-12-05         NaN
         1998-12-06         NaN
         1998-12-07         NaN
         1998-12-08    0.484211
         1998-12-09   -1.271930
         1998-12-10    0.467360
         Freq: D, dtype: float64
```

```
In [76]: ts.shift(-2)
```

```
Out[76]: 1998-12-05    0.467360
         1998-12-06    0.952550
         1998-12-07    0.371530
         1998-12-08   -1.454242
         1998-12-09         NaN
         1998-12-10         NaN
         Freq: D, dtype: float64
```

```
In [77]: ts.shift(1, freq='90S')
```

```
Out[77]: 1998-12-05 00:01:30    0.484211
         1998-12-06 00:01:30   -1.271930
         1998-12-07 00:01:30    0.467360
         1998-12-08 00:01:30    0.952550
         1998-12-09 00:01:30    0.371530
         1998-12-10 00:01:30   -1.454242
         Freq: D, dtype: float64
```



# Time Zone Handling

working with time zones is considered to be a pain. As a result many time series used choose to with with UTC, the succesor to GMT. IT is the current international standard. Time zones are expressed as offsets from UTC. In python, time zone information comes from pytz lrary. You can install it with pip or conda. pandas wraps pytz's functionality so you can ignore its API outside of the time zone names. Time zone names can be found interactively and in the docs.

```
In [78]: import pytz
pytz.common_timezones[-5:]
```

```
Out[78]: ['US/Eastern', 'US/Hawaii', 'US/Mountain', 'US/Pacific', 'UTC']
```

```
In [79]: # To get a time zone object from pytz, use pytz.timezone
tz = pytz.timezone('America/New_York')
tz
```

```
Out[79]: <DstTzInfo 'America/New_York' LMT-1 day, 19:04:00 STD>
```

## Time Zone Localization and Conversion

By default, Time series in pandas are time zone naive. For example, consider the following time series

```
In [80]: rng = pd.date_range('3/9/2012 9:20', periods=7, freq='D')
rng
```

```
Out[80]: DatetimeIndex(['2012-03-09 09:20:00', '2012-03-10 09:20:00',
                        '2012-03-11 09:20:00', '2012-03-12 09:20:00',
                        '2012-03-13 09:20:00', '2012-03-14 09:20:00',
                        '2012-03-15 09:20:00'],
                        dtype='datetime64[ns]', freq='D')
```

```
In [81]: ts = pd.Series(np.random.randn(len(rng)), index=rng)
ts
```

```
Out[81]: 2012-03-09 09:20:00    -2.070686
2012-03-10 09:20:00     1.916341
2012-03-11 09:20:00    -0.393583
2012-03-12 09:20:00     0.893082
2012-03-13 09:20:00     1.851536
2012-03-14 09:20:00    -0.172357
2012-03-15 09:20:00    -0.370833
Freq: D, dtype: float64
```

```
In [82]: print(ts.index.tz) # Shows the series does not have an associated time zone
None
```

```
In [83]: pd.date_range('3/9/2012 9:30', periods=10, freq='D', tz='UTC') #can pass a time zone
```

```
Out[83]: DatetimeIndex(['2012-03-09 09:30:00+00:00', '2012-03-10 09:30:00+00:00',
                        '2012-03-11 09:30:00+00:00', '2012-03-12 09:30:00+00:00',
                        '2012-03-13 09:30:00+00:00', '2012-03-14 09:30:00+00:00',
                        '2012-03-15 09:30:00+00:00', '2012-03-16 09:30:00+00:00',
                        '2012-03-17 09:30:00+00:00', '2012-03-18 09:30:00+00:00'],
                        dtype='datetime64[ns, UTC]', freq='D')
```

```
In [84]: ts_utc = ts.tz_localize('UTC') #localises time series to UTC
ts_utc
```

```
Out[84]: 2012-03-09 09:20:00+00:00    -2.070686
2012-03-10 09:20:00+00:00     1.916341
2012-03-11 09:20:00+00:00    -0.393583
2012-03-12 09:20:00+00:00     0.893082
2012-03-13 09:20:00+00:00     1.851536
2012-03-14 09:20:00+00:00    -0.172357
2012-03-15 09:20:00+00:00    -0.370833
Freq: D, dtype: float64
```

```
In [85]: ts_utc.index
```

```
Out[85]: DatetimeIndex(['2012-03-09 09:20:00+00:00', '2012-03-10 09:20:00+00:00',
                        '2012-03-11 09:20:00+00:00', '2012-03-12 09:20:00+00:00',
                        '2012-03-13 09:20:00+00:00', '2012-03-14 09:20:00+00:00',
                        '2012-03-15 09:20:00+00:00'],
                        dtype='datetime64[ns, UTC]', freq='D')
```

```
In [86]: ts_utc.tz_convert('America/New_York') # converts it to another time zone
```

```
Out[86]: 2012-03-09 04:20:00-05:00    -2.070686
2012-03-10 04:20:00-05:00     1.916341
2012-03-11 05:20:00-04:00    -0.393583
2012-03-12 05:20:00-04:00     0.893082
2012-03-13 05:20:00-04:00     1.851536
2012-03-14 05:20:00-04:00    -0.172357
2012-03-15 05:20:00-04:00    -0.370833
Freq: D, dtype: float64
```

```
In [87]: ts_eastern = ts.tz_localize('America/New_York') #localise to America/1
ts_eastern.tz_convert('UTC') #Converts to UTC
```

```
Out[87]: 2012-03-09 14:20:00+00:00    -2.070686
2012-03-10 14:20:00+00:00     1.916341
2012-03-11 13:20:00+00:00    -0.393583
2012-03-12 13:20:00+00:00     0.893082
2012-03-13 13:20:00+00:00     1.851536
2012-03-14 13:20:00+00:00    -0.172357
2012-03-15 13:20:00+00:00    -0.370833
Freq: D, dtype: float64
```

```
In [88]: ts_eastern.tz_convert('Europe/Berlin')
```

```
Out[88]: 2012-03-09 15:20:00+01:00    -2.070686
2012-03-10 15:20:00+01:00     1.916341
2012-03-11 14:20:00+01:00    -0.393583
2012-03-12 14:20:00+01:00     0.893082
2012-03-13 14:20:00+01:00     1.851536
2012-03-14 14:20:00+01:00    -0.172357
2012-03-15 14:20:00+01:00    -0.370833
Freq: D, dtype: float64
```

## Operations Between Different Time Zones

If two Series with different time zones are combined, the result will be UTC. Since the timestamps are stored under the hood in UTC.

```
In [89]: rng = pd.date_range('3/7/2012 9:30', periods=10, freq='B')
ts = pd.Series(np.random.randn(len(rng)), index=rng)
ts
```

```
Out[89]: 2012-03-07 09:30:00     0.605051
2012-03-08 09:30:00     0.481053
2012-03-09 09:30:00    -0.330596
2012-03-12 09:30:00     0.510970
2012-03-13 09:30:00    -0.340682
2012-03-14 09:30:00    -0.271859
2012-03-15 09:30:00     1.490435
2012-03-16 09:30:00    -1.887412
2012-03-19 09:30:00    -0.395401
2012-03-20 09:30:00     0.802360
Freq: B, dtype: float64
```

```
In [90]: ts1 = ts[:7].tz_localize('Europe/London')
         ts2 = ts1[2:].tz_convert('Europe/Moscow')
         result = ts1 + ts2
         result.index # Resultant time zone is UTC
```

```
Out[90]: DatetimeIndex(['2012-03-07 09:30:00+00:00', '2012-03-08 09:30:00+00:00',
                        '2012-03-09 09:30:00+00:00', '2012-03-12 09:30:00+00:00',
                        '2012-03-13 09:30:00+00:00', '2012-03-14 09:30:00+00:00',
                        '2012-03-15 09:30:00+00:00'],
                        dtype='datetime64[ns, UTC]', freq='B')
```

## Periods and Period Arithmetic

Periods represent timespanes, like days, months, quaters or years. The period is represented by the Period class

```
In [91]: p = pd.Period(2007, freq='A-Dec')
         p
```

```
Out[91]: Period('2007', 'A-DEC')
```

```
In [92]: # addidng and subtracting shift the periods date
         p + 5
```

```
Out[92]: Period('2012', 'A-DEC')
```

```
In [93]: # regular ranges or periods can be constructed with period_range
         rng = pd.period_range('2000-01-01', '2000-06-30', freq='M')
         rng
```

```
Out[93]: PeriodIndex(['2000-01', '2000-02', '2000-03', '2000-04', '2000-05',
                      '2000-06'], dtype='period[M]', freq='M')
```

```
In [94]: pd.Series(np.random.randn(len(rng)), index=rng) #can use as an axis in
```

```
Out[94]: 2000-01    0.651785
         2000-02    0.750168
         2000-03   -0.866046
         2000-04    0.470224
         2000-05   -0.949654
         2000-06   -1.326554
         Freq: M, dtype: float64
```

## period Frequency Conversion

for example should you have a period index of a year and wish to convert it into months

```
In [95]: p = pd.Period('2007', freq='A-DEC')
p
```

```
Out[95]: Period('2007', 'A-DEC')
```

```
In [96]: p.asfreq('M', how='start')
```

```
Out[96]: Period('2007-01', 'M')
```

```
In [97]: p.asfreq('M', how='end')
```

```
Out[97]: Period('2007-12', 'M')
```

PeriodIndex objects or time series can be similarly converted with the same semantics

```
In [98]: rng = pd.period_range('2006', '2012', freq='A-DEC')
p = pd.Series(np.random.randn(len(rng)), index=rng)
p
```

```
Out[98]: 2006    -1.061012
2007     0.163543
2008     0.666401
2009     2.197572
2010     1.363657
2011    -0.417862
2012    -0.852959
Freq: A-DEC, dtype: float64
```

```
In [99]: p.asfreq('B', how='end') # last business day of each year
```

```
Out[99]: 2006-12-29    -1.061012
2007-12-31     0.163543
2008-12-31     0.666401
2009-12-31     2.197572
2010-12-31     1.363657
2011-12-30    -0.417862
2012-12-31    -0.852959
Freq: B, dtype: float64
```

## Converting Timestamps to Periods (and Back)

Series and DataFrame objects indexed by timestamps can be converted to periods with the `to_period` method

```
In [100]: rng = pd.date_range('2000-01-01', periods=3, freq='M')
          ts = pd.Series(np.random.randn(3), index=rng)
          ts
```

```
Out[100]: 2000-01-31    1.408165
          2000-02-29    0.693272
          2000-03-31    0.189881
          Freq: M, dtype: float64
```

```
In [101]: pts = ts.to_period()
          pts
```

```
Out[101]: 2000-01    1.408165
          2000-02    0.693272
          2000-03    0.189881
          Freq: M, dtype: float64
```

Since periods refer to non-overlapping timespans, a timestamp can only belong to a single period. There however is no issue with having multiple numbers of the same period.

```
In [102]: rng = pd.date_range('1/29/2000', periods=6, freq='D')
          ts2 = pd.Series(np.random.randn(6), index=rng)
          ts2
```

```
Out[102]: 2000-01-29   -0.695630
          2000-01-30   -1.227850
          2000-01-31    0.655502
          2000-02-01   -0.434818
          2000-02-02    0.157314
          2000-02-03    0.174238
          Freq: D, dtype: float64
```

```
In [103]: pts2 = ts2.to_period('M')
          pts2
```

```
Out[103]: 2000-01   -0.695630
          2000-01   -1.227850
          2000-01    0.655502
          2000-02   -0.434818
          2000-02    0.157314
          2000-02    0.174238
          Freq: M, dtype: float64
```

```
In [104]: # To convert back to timestamp use to_timestamp()
```

```
In [105]: ts3 = pts2.to_timestamp(how='end')
          ts3
```

```
Out[105]: 2000-01-31 23:59:59.999999999    -0.695630
          2000-01-31 23:59:59.999999999    -1.227850
          2000-01-31 23:59:59.999999999     0.655502
          2000-02-29 23:59:59.999999999   -0.434818
          2000-02-29 23:59:59.999999999     0.157314
          2000-02-29 23:59:59.999999999     0.174238
          dtype: float64
```

## Resampling and Frequency conversion

resampling refers to the process of converting a time series from one frequency to another. downsampling, converts to lower frequency where upsampling refers to higher frequency. the resample method is the main workhorse for all frequency conversion.

```
In [106]: rng = pd.date_range('2000-01-01', periods=100, freq='D')
          ts = pd.Series(np.random.randn(len(rng)), index=rng)
          ts.head(8)
```

```
Out[106]: 2000-01-01    -1.066628
          2000-01-02    -1.394048
          2000-01-03     0.161837
          2000-01-04    -1.373027
          2000-01-05     0.055547
          2000-01-06    -1.287104
          2000-01-07    -0.352793
          2000-01-08     1.503156
          Freq: D, dtype: float64
```

```
In [107]: ts.resample('M').mean()
```

```
Out[107]: 2000-01-31    -0.400625
          2000-02-29    -0.156533
          2000-03-31     0.042490
          2000-04-30     0.069311
          Freq: M, dtype: float64
```

```
In [108]: ts.resample('M', kind='period').mean()
```

```
Out[108]: 2000-01    -0.400625
          2000-02    -0.156533
          2000-03     0.042490
          2000-04     0.069311
          Freq: M, dtype: float64
```

## Resample method arguments

freq -- String or DateOffset indicating desired resampled frequency axis -- Axis to resample on; default=0  
 fill\_method -- How to interpolate when upsampling as in 'ffill', or 'bfill';  
 closed -- In downsampling, which end of each interval is closed 'right' or 'left'.  
 label -- In downsampling, how to label the aggregated result with the 'right' or 'left' bin edge  
 limit -- When forward or backward filling, the maximum number of periods to fill.  
 kind -- Aggregate to periods or timestamps 'period' or 'timestamp'  
 convention -- When resampling periods, the convention('start' or 'end' for converting the low-frequency period to high frequency'; defaults to 'start'.

## Downsampling

when downsampling you need to think about

- Which side of each interval is closed
- How to label each aggregated bin, either with the start of the interval or the end.

```
In [110]: rng = pd.date_range('2000-01-01', periods=12, freq='T')
          ts = pd.Series(np.arange(len(rng)), index=rng)
          ts
```

```
Out[110]: 2000-01-01 00:00:00    0
          2000-01-01 00:01:00    1
          2000-01-01 00:02:00    2
          2000-01-01 00:03:00    3
          2000-01-01 00:04:00    4
          2000-01-01 00:05:00    5
          2000-01-01 00:06:00    6
          2000-01-01 00:07:00    7
          2000-01-01 00:08:00    8
          2000-01-01 00:09:00    9
          2000-01-01 00:10:00   10
          2000-01-01 00:11:00   11
          Freq: T, dtype: int64
```

```
In [111]: ts.resample('5min', closed='right').sum()
```

```
Out[111]: 1999-12-31 23:55:00    0
          2000-01-01 00:00:00   15
          2000-01-01 00:05:00   40
          2000-01-01 00:10:00   11
          Freq: 5T, dtype: int64
```



```
In [113]: ts.resample('5min', closed='right', label='right').sum() # changes the
```

```
Out[113]: 2000-01-01 00:00:00    0
          2000-01-01 00:05:00    15
          2000-01-01 00:10:00    40
          2000-01-01 00:15:00    11
          Freq: 5T, dtype: int64
```

## Open-High-Low-Close resampling

in finance a popular way to aggregate a time series is to compute four values for each bucket. By using ohlc aggregate function you will obtain a DataFrame having columns containing these four aggregates

```
In [114]: ts.resample('5min').ohlc()
```

```
Out[114]:
```

	open	high	low	close
2000-01-01 00:00:00	0	4	0	4
2000-01-01 00:05:00	5	9	5	9
2000-01-01 00:10:00	10	11	10	11

## Upsampling and Interpolation

when converting from a low frequency to a higher frequency no aggregation is needed.

```
In [118]: frame = pd.DataFrame(np.random.randn(2,4), index=pd.date_range('1/1/2000', periods=2))
          frame
```

```
Out[118]:
```

	Colorado	Texas	New York	Ohio
2000-01-05	1.312761	-0.186424	1.127892	-0.630408
2000-01-12	-0.880612	-0.578471	0.799949	1.848771

```
In [120]: df_daily = frame.resample('D').asfreq()
df_daily
```

Out[120]:

	Colorado	Texas	New York	Ohio
2000-01-05	1.312761	-0.186424	1.127892	-0.630408
2000-01-06	NaN	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN	NaN
2000-01-08	NaN	NaN	NaN	NaN
2000-01-09	NaN	NaN	NaN	NaN
2000-01-10	NaN	NaN	NaN	NaN
2000-01-11	NaN	NaN	NaN	NaN
2000-01-12	-0.880612	-0.578471	0.799949	1.848771

```
In [123]: frame.resample('D').ffill() # fill carries the data foward and resamp
# as ffill reduces the time frequency and therefore rows
```

Out[123]:

	Colorado	Texas	New York	Ohio
2000-01-05	1.312761	-0.186424	1.127892	-0.630408
2000-01-06	1.312761	-0.186424	1.127892	-0.630408
2000-01-07	1.312761	-0.186424	1.127892	-0.630408
2000-01-08	1.312761	-0.186424	1.127892	-0.630408
2000-01-09	1.312761	-0.186424	1.127892	-0.630408
2000-01-10	1.312761	-0.186424	1.127892	-0.630408
2000-01-11	1.312761	-0.186424	1.127892	-0.630408
2000-01-12	-0.880612	-0.578471	0.799949	1.848771

## Resampling with Periods

Resampling data indexed by periods is similar to timestamps

```
In [125]: frame = pd.DataFrame(np.random.randn(24, 4), index=pd.period_range('1-
frame.head(3)
```

Out[125]:

	Colorado	Texas	New York	Ohio
<b>2000-01</b>	-0.952643	-0.856533	0.024391	0.108535
<b>2000-02</b>	0.446466	-0.460336	-2.619437	-0.252802
<b>2000-03</b>	0.735194	-0.892583	1.422746	0.216618

```
In [126]: annual_frame = frame.resample('A-Dec').mean()
annual_frame
```

Out[126]:

	Colorado	Texas	New York	Ohio
<b>2000</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001</b>	0.472907	0.223241	0.018972	-0.343108

```
In [127]: # Upsampling is more nuanced, as you must make a decision about which
# to place the values before resampling
# Q-DEC: quarterly, year ending in December
annual_frame.resample('Q-DEC').ffill()
```

Out[127]:

	Colorado	Texas	New York	Ohio
<b>2000Q1</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2000Q2</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2000Q3</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2000Q4</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001Q1</b>	0.472907	0.223241	0.018972	-0.343108
<b>2001Q2</b>	0.472907	0.223241	0.018972	-0.343108
<b>2001Q3</b>	0.472907	0.223241	0.018972	-0.343108
<b>2001Q4</b>	0.472907	0.223241	0.018972	-0.343108

```
In [129]: annual_frame.resample('Q-DEC', convention='end').ffill()
```

Out[129]:

	Colorado	Texas	New York	Ohio
<b>2000Q4</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001Q1</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001Q2</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001Q3</b>	0.164265	-0.109949	0.040561	-0.008732
<b>2001Q4</b>	0.472907	0.223241	0.018972	-0.343108

**In downsampling, the target frequency must be a subperiod of the source frequency.**

**In upsampling, the target frequency must be a superperiod fo the source frequency.**

In [ ]: